

Computer Vision

Assignment 12

Alireza Moradi

December 18, 2020

1 Shape Descriptors

$$Compactness = \frac{4\pi \times Area}{Perimeter^2}$$

$$Solidity = \frac{Area}{ConvexArea}$$

$$Eccentricity = \frac{MinorAxisLength}{MajorAxisLength}$$

1.1 Rhombus


- Compactness: $\frac{4\pi \times a^2}{20a^2} = \frac{1}{5}\pi$
- Solidity: 1 (because lozenge is convex)
- Eccentricity: $\frac{a}{2a} = \frac{1}{2}$

1.2 Rectangle

- Compactness: $\frac{4\pi \times 2a^2}{36a^2} = \frac{2}{9}\pi$
- Solidity: 1 (because rectangle is convex)
- Eccentricity: $\frac{a\frac{\sqrt{5}}{2}}{a\sqrt{5}} = \frac{1}{2}$

2

$$\begin{aligned} 7 &= (00000111)_2 \xrightarrow{Rotate90} (11000001)_2 = 193 \\ 62 &= (00111110)_2 \xrightarrow{Rotate90} (10001111)_2 = 143 \\ 112 &= (01110000)_2 \xrightarrow{Rotate90} (00011100)_2 = 28 \\ 255 &= (11111111)_2 \xrightarrow{Rotate90} (11111111)_2 = 255 \end{aligned}$$



3 Implementations

3.1 SVM

3.1.1 HOG

I used `skimage.feature.hog` and got results with 8 different orientations and 8×8 pixels per cell and 3×3 cells per block and this resulted in a classification with 93.7% accuracy.




Figure 1: Confusion Matrix

```
1 def classify_hog(x_train, y_train, x_test):
2     '''
3     Classifies images with HOG.
4
5     Parameters:
6         x_train(numpy.ndarray) : train data
7         y_train(numpy.ndarray) : train labels
8         x_test(numpy.ndarray) : test data
9
10    outputs:
11        predicted_labels(numpy.ndarray): labels that predicted
12    '''
13
14    x_t = x_train.copy()
15    y_t = y_train.copy()
16    x_te = x_test.copy()
17    predicted_labels = None
```

```
18
19     hog_image = np.array([hog(x, orientations=8, pixels_per_cell=(8, 8),
20                             cells_per_block=(3, 3), visualize=True,
21                             multichannel=False)[0] for x in x_train])
22
23     hog_image_test = np.array([hog(x, orientations=8, pixels_per_cell=(8, 8),
24                                   cells_per_block=(3, 3), visualize=True,
25                                   multichannel=False)[0] for x in x_test])
26
27     svm = LinearSVC()
28     svm.fit(hog_image, y_train)
29     predicted_labels = svm.predict(hog_image_test)
30     return predicted_labels
```

3.1.2 Shape Descriptors

I computed the contours with opencv and chose the largest one. Using the largest contour I computed

- **Extreme Points**: topmost, bottommost, rightmost and leftmost points of the object.
- **Aspect Ratio**: ratio of width to height of bounding rect of the object.
- **Solidity**: the ratio of contour area to its convex hull area.

of each data and used these to do the classification which resulted in 59.2% accuracy.




Figure 2: Confusion Matrix

```

1 def extract_shape_desc(x):
2     _,thresh = cv2.threshold(x,170,255,cv2.THRESH_BINARY)
3     contours,hierarchy = cv2.findContours(thresh,cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_NONE)
4     contour_sizes = [(cv2.contourArea(contour), contour) for contour in contours]
5     biggest_contour = max(contour_sizes, key=lambda x: x[0])[1]
6     cnt = biggest_contour
7
8     leftmost = np.array(cnt[cnt[:, :, 0].argmin()][0]).reshape(2,1)
9     rightmost = np.array(cnt[cnt[:, :, 0].argmax()][0]).reshape(2,1)
10    topmost = np.array(cnt[cnt[:, :, 1].argmin()][0]).reshape(2,1)
11    bottommost = np.array(cnt[cnt[:, :, 1].argmax()][0]).reshape(2,1)
12
13    area = cv2.contourArea(cnt)
14    hull = cv2.convexHull(cnt)

```

```

15 hull_area = cv2.contourArea(hull)
16 solidity = float(area)/(hull_area+1e-7)
17
18 x,y,w,h = cv2.boundingRect(cnt)
19 aspect_ratio = float(w)/h
20
21 return np.concatenate((leftmost,rightmost,
22                        topmost,bottommost,[[solidity],[aspect_ratio]]),axis=0)

```

```

1 def classify_shape_desc(x_train, y_train, x_test):
2     '''
3     Classifies images by using shape descriptors.
4
5     Parameters:
6         x_train(numpy.ndarray) : train data
7         y_train(numpy.ndarray) : train labels
8         x_test(numpy.ndarray) : test data
9
10    outputs:
11        predicted_labels(numpy.ndarray): labels that predicted
12    '''
13
14    x_t = x_train.copy()
15    y_t = y_train.copy()
16    x_te = x_test.copy()
17    predicted_labels = None
18
19    data = np.squeeze([extract_shape_desc(x) for x in x_train])
20    test = np.squeeze([extract_shape_desc(x) for x in x_test])
21
22    svm = LinearSVC()
23    svm.fit(data,y_train)
24    predicted_labels = svm.predict(test)
25
26    return predicted_labels

```

3.1.3 Manual LBP

I computed LBP manually (which takes a lot of time because of those nested for loops) and computed the normalized histogram of LBP with 256 bins from 0 to 255. this resulted in 70.6% accuracy.

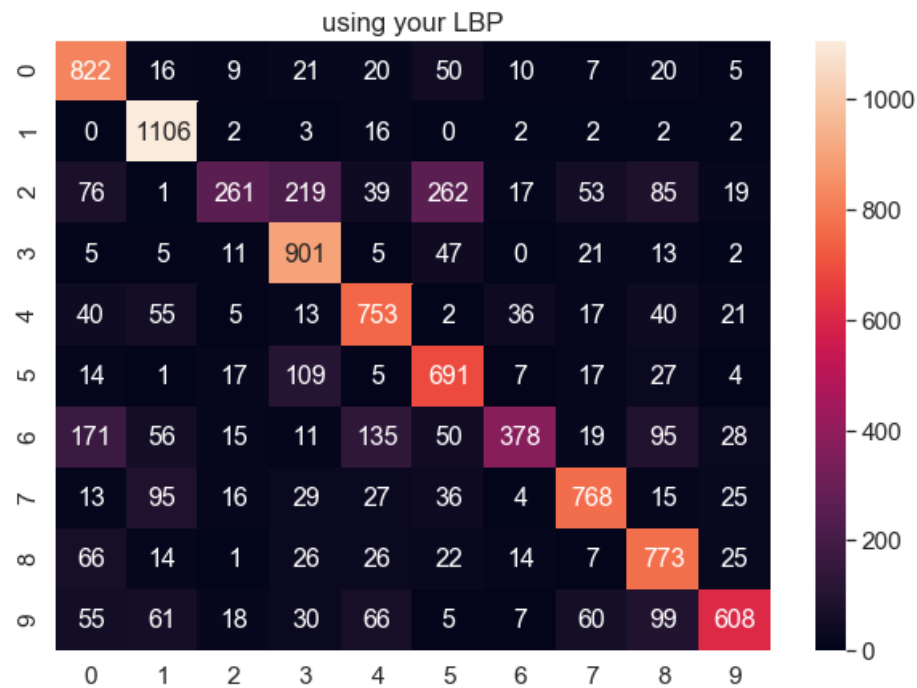


Figure 3: Confusion Matrix

[Source 1](#) [Source 2](#)

```

1 def LBP(img):
2     '''
3     Extracts LBP features from the input image.
4
5     Parameters:
6         img(numpy.ndarray) : image data
7     outputs:
8         output: LBP features
9     '''
10    R,C = img.shape
11
12    transformed = np.zeros((26, 26))
13    for i in range (1,R-1):
14        for k in range (1,C-1):
15
16            binary_matrix = np.zeros((3, 3))
17            sub_mat = img[i-1:i+2,k-1:k+2]
18            mid = sub_mat[1,1]
19            higher = np.where(sub_mat > mid)

```

```

20         lower = np.where(sub_mat <= mid)
21         binary_matrix[higher] = 1
22         binary_matrix[lower] = 0
23         transformed[i-1,k-1] = np.sum(np.multiply(multip_matrix,binary_matrix))
24
25     hist, _ = np.histogram(transformed, bins=256, density=True, range=(0, 256))
26
27     return hist

```

```

1 def classify_your_lbp(x_train, y_train, x_test):
2     '''
3     Classifies images by using your LBP.
4
5     Parameters:
6         x_train(numpy.ndarray) : train data
7         y_train(numpy.ndarray) : train labels
8         x_test(numpy.ndarray) : test data
9
10    outputs:
11        predicted_labels(numpy.ndarray): labels that predicted
12    '''
13    data = [LBP(x) for x in x_train]
14    test = [LBP(x) for x in x_test]
15
16    data = np.array(data)
17    test = np.array(test)
18    svm = LinearSVC()
19    svm.fit(data,y_train)
20    predicted_labels = svm.predict(test)
21
22    return predicted_labels

```


3.1.4 LBP

I computed LBP by using `skimage.feature.local_binary_pattern` and computed normalized histogram. This was much faster but resulted in 57% accuracy.

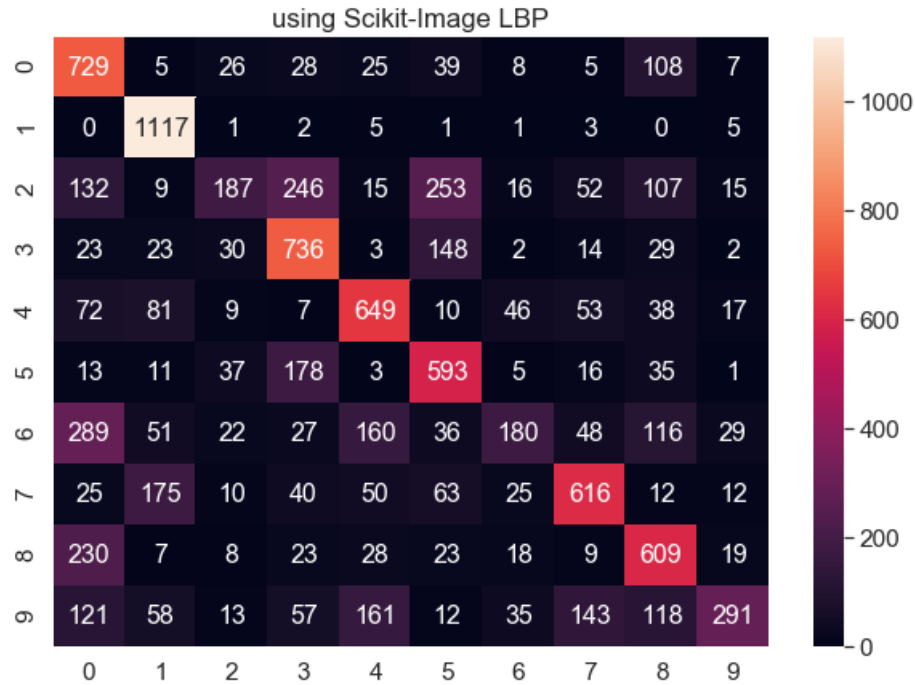


Figure 4: Confusion Matrix

```
1 def classify_skimage_lbp(x_train, y_train, x_test):
2     '''
3     Classifies images by using Scikit-Image LBP.
4
5     Parameters:
6         x_train(numpy.ndarray) : train data
7         y_train(numpy.ndarray) : train labels
8         x_test(numpy.ndarray) : test data
9
10    outputs:
11        predicted_labels(numpy.ndarray): labels that predicted
12    '''
13
14    data = [np.histogram(local_binary_pattern(x,8,1), density=True
15                      , bins=256, range=(0, 256))[0] for x in x_train]
16    test = [np.histogram(local_binary_pattern(x,8,1), density=True
17                      , bins=256, range=(0, 256))[0] for x in x_test]
18    data = np.array(data)
19    test = np.array(test)
20    svm = LinearSVC()
21    svm.fit(data,y_train)
```

```
22     predicted_labels = svm.predict(test)
23
24     return predicted_labels
```

3.2 Binary Classification

I used the following features (which are all scale invariant):

- **Extent**: the ratio of contour area to bounding rectangle area.
- **Aspect Ratio**: ratio of width to height of bounding rect of the object.
- **Solidity**: the ratio of contour area to its convex hull area.

Observing and playing with the numbers from the above features a little, I managed to classify all the leaves correctly.

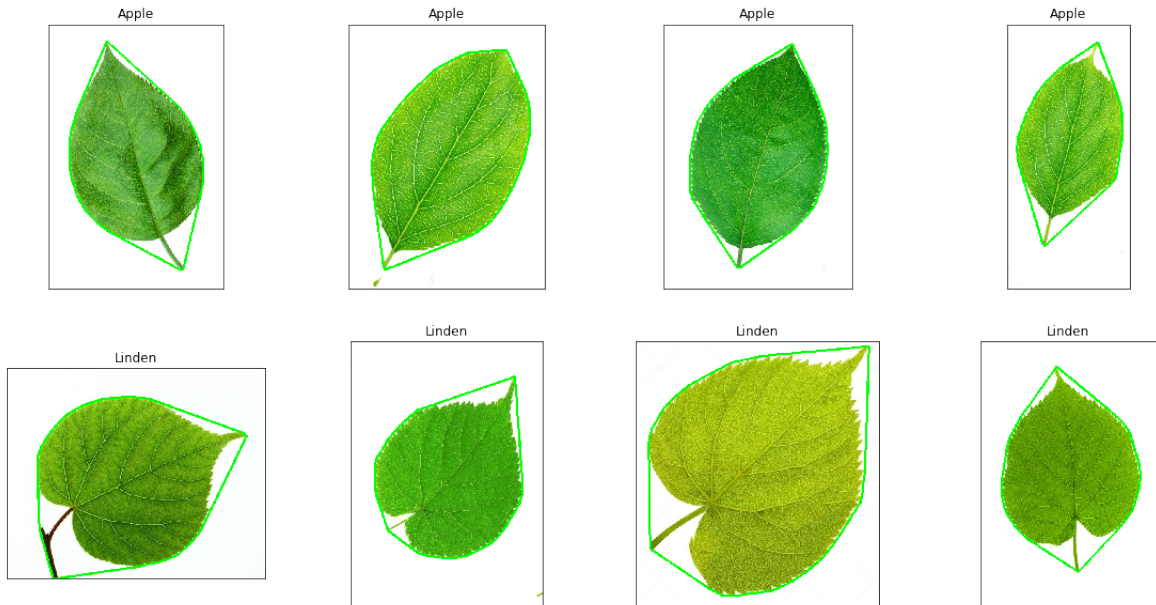


Figure 5: Classified leaves and their convex hull

```
1 def classify_leaf(image):
2     '''
3     Classifies the input image to only two classes of leaves.
4
5     Parameters:
6         image (numpy.ndarray): The input image.
7
8     Returns:
9         int: The class of the image. 1 == apple, 0 == linden
10    '''
11    # apple = 0
12    # linden = 1
13    leaf_type = 0
14
15    #Write your code here
16    gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
```

```

17     ret,thresh = cv2.threshold(gray,220,255,cv2.THRESH_BINARY_INV)
18     contours,hierarchy = cv2.findContours(thresh,cv2.RETR_TREE, cv2.CHAIN_APPROX_SIMPLE)
19
20     contour_sizes = [(cv2.contourArea(contour), contour) for contour in contours]
21     biggest_contour = max(contour_sizes, key=lambda x: x[0])[1]
22     cnt = biggest_contour
23
24     area = cv2.contourArea(cnt)
25     hull = cv2.convexHull(cnt)
26     hull_area = cv2.contourArea(hull)
27     solidity = float(area)/(hull_area)
28
29     x,y,w,h = cv2.boundingRect(cnt)
30     aspect_ratio = float(w)/h
31
32     area = cv2.contourArea(cnt)
33     x,y,w,h = cv2.boundingRect(cnt)
34     rect_area = w*h
35     extent = float(area)/rect_area
36
37     cv2.drawContours(image, [hull], -1, (0,255,0), 3)
38
39     if solidity <= 0.875:
40         if aspect_ratio >= 0.75:
41             return 0
42         else:
43             if extent <= 0.6:
44                 return 0
45             else:
46                 return 1
47     else:
48         if aspect_ratio <= 0.75:
49             return 1
50         else:
51             if extent <= 0.6:
52                 return 1
53             else:
54                 return 0
55
56     return leaf_type

```